

Part 1: Mutual Information Classification

```
# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.impute import SimpleImputer
from sklearn.feature_selection import SelectKBest, mutual_info_classif

# Load the dataset (replace 'your_dataset.csv' with the actual file path)
df = pd.read_csv('loan.csv')

# Display the first few rows of the dataset
print("Original Dataset:")
print(df.head())
```

Original Dataset:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

```
# Handle null values (replace 'most_frequent' with an appropriate imputation strategy)
imputer = SimpleImputer(strategy='most_frequent')
df = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)

# Display the dataset after handling null values
print("\nDataset after null value handling:")
print(df.head())
```

Dataset after null value handling:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	120.0	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

Separate features (X) and target variable (y)

```
X = df.drop('Loan_Status', axis=1)
```

```
y = df['Loan_Status']
```

Label encoding for categorical variables

```
label_encoder = LabelEncoder()
```

```
X_encoded = X.apply(label_encoder.fit_transform)
```

Display the dataset after label encoding

```
print("\nDataset after label encoding:")
```

```
print(X_encoded.head())
```

Dataset after label encoding:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	0	1	0	0	0	0	
1	1	1	1	1	0	0	
2	2	1	1	0	0	1	
3	3	1	1	0	1	0	
4	4	1	0	0	0	0	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	376	0	73	8	
1	306	60	81	8	
2	139	0	26	8	
3	90	160	73	8	
4	381	0	94	8	

	Credit_History	Property_Area
0	1	2
1	1	0
2	1	2
3	1	2
4	1	2

```
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y,
test_size=0.2, random_state=42)

# Use SelectKBest to select the top K features based on mutual
information scores
k_best_selector = SelectKBest(score_func=mutual_info_classif, k=3) #
Choose an appropriate value of K
X_train_kbest = k_best_selector.fit_transform(X_train, y_train)

# Get the indices of the selected features
selected_feature_indices = k_best_selector.get_support(indices=True)

# Print the names or indices of the selected features
selected_feature_names = X_train.columns[selected_feature_indices]
print("\nSelected Features:")
print(selected_feature_names)
```

```
Selected Features:
Index(['Self_Employed', 'Credit_History', 'Property_Area'],
dtype='object')
```

Part 2: Mutual Information Regression

```
# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.impute import SimpleImputer
from sklearn.feature_selection import SelectKBest,
mutual_info_regression
from sklearn.ensemble import RandomForestRegressor

# Load the dataset (replace 'housing_dataset.csv' with the actual file
path)
df = pd.read_csv('housing.csv')

# Display the first few rows of the dataset
print("Original Dataset:")
print(df.head())
```

Original Dataset:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0		
2										
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0		
5										
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0		
9										
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0		
2										
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0		
12										

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

```
# Handle null values (replace 'mean' with an appropriate imputation strategy)
```

```
imputer = SimpleImputer(strategy='most_frequent')
```

```
df = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)
```

```
# Display the dataset after null value handling
```

```
print("\nDataset after null value handling:")
```

```
print(df.head())
```

Dataset after null value handling:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	Grvl	Reg	
1	2	20	RL	80.0	9600	Pave	Grvl	Reg	
2	3	60	RL	68.0	11250	Pave	Grvl	IR1	

3	4	70	RL	60.0	9550	Pave	Grvl	IR1
4	5	60	RL	84.0	14260	Pave	Grvl	IR1

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature
MiscVal \							
0	Lvl	AllPub	...	0	Gd	MnPrv	Shed
0							
1	Lvl	AllPub	...	0	Gd	MnPrv	Shed
0							
2	Lvl	AllPub	...	0	Gd	MnPrv	Shed
0							
3	Lvl	AllPub	...	0	Gd	MnPrv	Shed
0							
4	Lvl	AllPub	...	0	Gd	MnPrv	Shed
0							

	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	2	2008	WD	Normal	208500
1	5	2007	WD	Normal	181500
2	9	2008	WD	Normal	223500
3	2	2006	WD	Abnorml	140000
4	12	2008	WD	Normal	250000

[5 rows x 81 columns]

Separate features (X) and target variable (y)

```
X = df.drop('SalePrice', axis=1)
```

```
y = df['SalePrice']
```

Label encoding for categorical variables

```
label_encoder = LabelEncoder()
```

Apply label encoding to each categorical column

```
for col in X.select_dtypes(include=['object']).columns:
    X[col] = label_encoder.fit_transform(X[col])
```

Display the dataset after label encoding

```
print("\nDataset after label encoding:")
```

```
print(X.head())
```

Dataset after label encoding:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley
LotShape \							
0	0	5	3	36	327	1	0
3							
1	1	0	3	51	498	1	0
3							
2	2	5	3	39	702	1	0
0							

3	3	6	3	31	489	1	0
0							
4	4	5	3	55	925	1	0
0							

	LandContour	Utilities	...	ScreenPorch	PoolArea	PoolQC	
Fence \							
0	3	0	...	0	0	2	2
1	3	0	...	0	0	2	2
2	3	0	...	0	0	2	2
3	3	0	...	0	0	2	2
4	3	0	...	0	0	2	2

	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition	
0	2	0	1	2	8		4
1	2	0	4	1	8		4
2	2	0	8	2	8		4
3	2	0	1	0	8		0
4	2	0	11	2	8		4

[5 rows x 80 columns]

Split the dataset into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

Use RandomForestRegressor to estimate mutual information scores

```
rf_regressor = RandomForestRegressor()
rf_regressor.fit(X_train, y_train)
mutual_info_scores = rf_regressor.feature_importances_
```

Use SelectKBest to select the top K features based on mutual information scores

```
k_best_selector = SelectKBest(score_func=mutual_info_regression, k=3)
```

Choose an appropriate value of K

```
X_train_kbest = k_best_selector.fit_transform(X_train, y_train)
```

Get the indices of the selected features

```
selected_feature_indices = k_best_selector.get_support(indices=True)
```

Print the names or indices of the selected features

```
selected_feature_names = X_train.columns[selected_feature_indices]
```

```
print("\nSelected Features:")
```

```
print(selected_feature_names)
```

```
Selected Features:
Index(['OverallQual', 'GrLivArea', 'GarageCars'], dtype='object')
```

Part 3 : Linear Regression on the Housing Dataset

```
# Import necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer

# Load the Housing dataset (replace 'housing_dataset.csv' with the
actual file path)
df = pd.read_csv('housing.csv')
```

```
# Display the first few rows of the dataset
print("Original Dataset:")
print(df.head())
```

Original Dataset:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape
0	1	60	RL	65.0	8450	Pave	NaN	Reg
1	2	20	RL	80.0	9600	Pave	NaN	Reg
2	3	60	RL	68.0	11250	Pave	NaN	IR1
3	4	70	RL	60.0	9550	Pave	NaN	IR1
4	5	60	RL	84.0	14260	Pave	NaN	IR1

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0
5	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0
9	Lvl	AllPub	...	0	NaN	NaN	NaN	0
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0

4	Lvl	AllPub	...	0	NaN	NaN	NaN	0
---	-----	--------	-----	---	-----	-----	-----	---

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

Handle null values (replace 'mean' with an appropriate imputation strategy)

imputer = SimpleImputer(strategy='most_frequent')

df = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)

Display the dataset after null value handling

print("\nDataset after null value handling:")

print(df.head())

Dataset after null value handling:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	Grvl	Reg	
1	2	20	RL	80.0	9600	Pave	Grvl	Reg	
2	3	60	RL	68.0	11250	Pave	Grvl	IR1	
3	4	70	RL	60.0	9550	Pave	Grvl	IR1	
4	5	60	RL	84.0	14260	Pave	Grvl	IR1	

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature
--	-------------	-----------	-----	----------	--------	-------	-------------

MiscVal	\
0	Lvl AllPub ... 0 Gd MnPrv Shed
0	
1	Lvl AllPub ... 0 Gd MnPrv Shed
0	
2	Lvl AllPub ... 0 Gd MnPrv Shed
0	
3	Lvl AllPub ... 0 Gd MnPrv Shed
0	
4	Lvl AllPub ... 0 Gd MnPrv Shed
0	

	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	2	2008	WD	Normal	208500
1	5	2007	WD	Normal	181500
2	9	2008	WD	Normal	223500
3	2	2006	WD	Abnorml	140000
4	12	2008	WD	Normal	250000


```

[5 rows x 81 columns]

# Separate features (X) and target variable (y)
X = df.drop('YrSold', axis=1)
y = df['YrSold']

# Label encoding for categorical variables
label_encoder = LabelEncoder()

# Apply label encoding to each categorical column
for col in X.select_dtypes(include=['object']).columns:
    X[col] = label_encoder.fit_transform(X[col])

# Split the dataset into training and testing sets (80:20 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Scale features using StandardScaler (optional but can improve model
performance)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Fit a linear regression model to the training data
model = LinearRegression()
model.fit(X_train_scaled, y_train)

# Predict house prices for the testing data
y_pred = model.predict(X_test_scaled)

# Evaluate the model's performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print the MSE and R^2 values
print(f"\nMean Squared Error (MSE): {mse:.2f}")
print(f"Coefficient of Determination (R^2): {r2:.2f}")

Mean Squared Error (MSE): 1.82
Coefficient of Determination (R^2): 0.00

# Plot a scatter plot between predicted and actual house prices
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='blue', alpha=0.6)
plt.title('Predicted vs Actual House Prices')
plt.xlabel('Actual House Prices')
plt.ylabel('Predicted House Prices')
plt.show()

```

